Voyager 1 — Data Flow

High-level view of data moving through each system

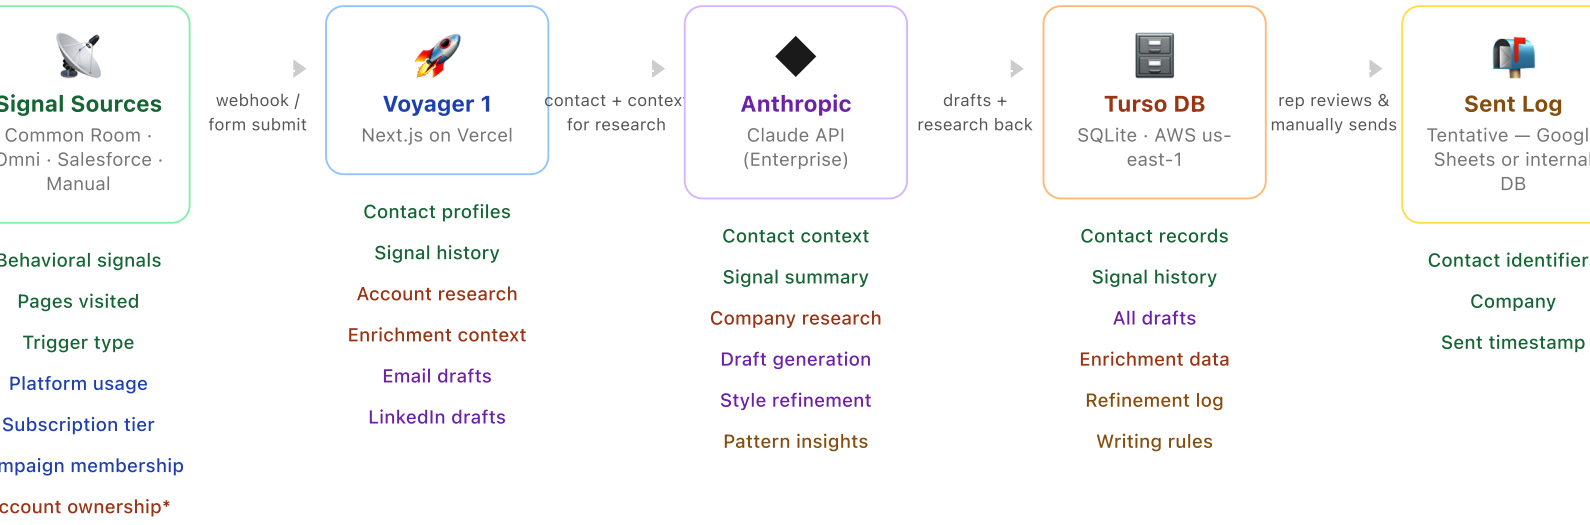

📡 Signal Sources

Common Room will push real-time behavioral signals. Omni will deliver daily platform sign-up data. Salesforce marketing campaigns will surface campaign members and engagement history. Reps will add leads manually. All routes will converge into Voyager 1.

* Account ownership mapping (Salesforce AE assignment, territory, open opportunities) will be a future integration.

🚀 Voyager 1

Voyager 1 will orchestrate the full workflow. It will queue leads for rep review, manage the draft lifecycle, and store enrichment context added by the rep (Common Room paste, Salesforce notes).

◆ Anthropic

Anthropic will receive contact and account context, research the company via web search, generate personalized email and LinkedIn drafts, and refine on rep feedback. The Enterprise contract will ensure no training data usage and no retention after request completion.

🗄️ Turso DB

Turso will serve as the persistent store for all contacts, accounts, drafts, and signal history. Every refinement the rep makes will be logged and fed back into future generations as style examples.

📬 Sent Log

When a rep clicks Send, outreach activity will be logged for pipeline tracking. Minimal data — contact identifiers, company, timestamp.

Tentative: destination TBD — Google Sheets via Apps Script or an internal DB. Internal DB would be preferable at scale.

Design principle: The rep will always review drafts before sending. Email send will never be automated. Voyager 1 will not have a direct integration into Gmail or any email client — the rep will copy the draft and send manually.

SECURITY ASSUMPTIONS & SAFEGUARDS

🔐 Access Control

- A single shared password (SITE_SECRET) will be required to access the app
- Sessions will be stored in an httpOnly, secure, sameSite=strict cookie — not accessible to JavaScript
- Middleware will enforce auth on every route and will fail closed if the secret is not configured
- Ingest webhooks will use a separate secret (INGEST_SECRET) validated on every

🔑 Secret Management

- All secrets will be stored in Vercel encrypted environment variables — never in source code or git
- Local development will use .env.local, which will be gitignored by default
- Database credentials (URL + auth token) will be server-side only — never exposed to the browser
- The sent log webhook URL will be server-side only — never sent to the client

🗄️ Database

- The database will require a signed JWT token for every connection — no unauthenticated access will be possible
- Data will be encrypted at rest
- Tokens will be immediately invalidatable and rotatable if compromised
- Lead deletion will be soft — records will be flagged rather than permanently removed, preserving an audit trail

incoming request

◆ AI Data Handling

- The Anthropic Enterprise contract will ensure contact data sent to the API will not be used for model training
- There will be no data retention on Anthropic's side after request completion
- All Anthropic API calls will be server-side — the API key will never be exposed to the browser

🌍 GDPR / Compliance

- The Anthropic Enterprise agreement will cover data processing for EMEA contacts
- A Vercel DPA will be required — available self-serve at vercel.com/legal/dpa
- A Turso DPA will be required — a DB migration to an alternative provider may be needed if unavailable
- Legitimate interests will serve as the legal basis for B2B outreach to contacts showing intent signals

⚠️ Known Gaps

- Single shared password — per-user accounts and action-level audit logging would be needed at larger team scale
- A hard-delete / right-to-erasure path will be needed to support GDPR erasure requests
- A cross-border transfer mechanism will need to be formalised for EMEA data given US-region database hosting
- Delete protection on the database will need to be enabled to prevent accidental data loss